

AI Critique

During HW05, Codex accelerated source tracing, JMeter generation, execution, and raw-log calculation, but it initially optimized for endpoint coverage instead of business coherence. The first FR03-FR09-FR16 chain technically exercised authentication, coupon, write, and read APIs, yet password recovery, coupon validation, and admin import did not form one user goal. The student's team-workflow comparison exposed that weakness, and the plan was replaced with WF-04: an admin logs in, imports a product, searches it, and opens the correlated detail record. AI also reused the previous run date in generated filenames until execution evidence forced a correction to 20260901. During the first WF-04 soak, Codex initially trusted a 600-second scheduler; the host clock jumped forward, leaving only 305.20 seconds of JTL traffic. Retaining that attempt as inconclusive and rerunning 410 iterations per thread produced 618.94 seconds and exactly 12,300 completed workflows. Task 2 showed further semantic risk: JMeter's 79.49 HTTP samples/s is not 79.49 workflows/s; the final-sampler count gives 19.87 workflows/s. A 1,151 ms maximum also does not override p95 4 ms and p99 6 ms, while a 136.67 MB RSS peak followed by 51.86 MB final RSS is not a memory ceiling. These errors came from treating configuration labels and aggregate summaries as conclusions, and from reusing generic database advice without checking query shape or SQLite architecture. The learned principle is evidence-gated collaboration: first validate one coherent business goal, then verify scope, timestamps, state deltas, stage boundaries, final-sampler completion, and source feasibility before accepting AI labels, thresholds, or recommendations.

Word count: 247.