

AI Critique

During HW05, Codex was useful for reading the assignment, tracing the real Express/SQLite implementation, generating JMeter plans, and checking raw JTL data, but its first outputs were not submission-ready. The initial generated XML scoped the CSV Data Set, headers, and think-time timer inside only the first thread group. A one-user smoke test passed, so that defect would have remained hidden in the later Stress and Spike stages without a structural XML audit. The first Load attempt also exposed an incomplete readiness assumption: Codex treated an accepting HTTP port as proof that the database was initialized, while `database.js` was still asynchronously executing destructive DROP/CREATE/seed statements. The performance accounts disappeared during the run, producing application errors that could easily have been misreported as load failures. Finally, the first soak proposal reused six virtual users without evidence; it was corrected to thirty only after a real thirty-user spike completed with zero errors. These failures came from optimistic defaults, insufficient attention to configuration scope, and confusing service availability with state readiness. They also show why a maximum response time or total HTTP sample rate must not be casually called an SLA or user-journey throughput: one complete workflow contains six HTTP samples, and an isolated outlier does not describe the percentile distribution. The main lesson is that collaboration with AI should be evidence-gated. Every generated plan needs a code-level scope review, a one-user business smoke test, a clean and stable data precondition, and comparison against full raw logs. AI can accelerate construction and analysis, but human ownership means challenging its assumptions before accepting its labels, thresholds, or conclusions.

Word count: 263.