

AI Critique

During HW05, Codex accelerated source tracing, JMeter generation, execution orchestration, and raw-log calculation, but its initial labels were more confident than the evidence. It first scoped shared JMeter configuration inside only one thread group; a one-user smoke passed even though later Stress and Spike groups would have missed data and timers. It then treated an open HTTP port as database readiness while asynchronous destructive initialization was still running. Task 2 exposed a different failure mode: generic performance language. The first analysis called 82.29 HTTP samples/s “workflows/s,” although the raw Soak log contains 8,207 successful final samplers, or 13.69 completed workflows/s over the overall window. It compared the complete Stress and Spike averages even though Spike includes two low-load 3-VU periods; the relevant burst was 82.96 samples/s. It also promoted a 222 ms maximum to a tail-latency problem despite p95 5 ms and p99 6 ms, and called a transient 140.47 MB RSS peak a memory ceiling even though RSS ended at 78.50 MB. Generic advice to add a coupon index and connection pool ignored that coupon code is already unique and the application uses one SQLite handle. These errors arose because aggregate summaries hide workload boundaries and language models reuse familiar optimization patterns without checking query shape or architecture. The principle learned is evidence-gated collaboration: retain the first AI output, recompute metrics from raw rows, compare equivalent stages, inspect the exact source path, and separate measured baselines from proposed regression gates and approved SLAs. AI may draft hypotheses quickly, but a human must own metric semantics, feasibility, and final acceptance.

Word count: 260.